

## practiceQuestions2(Answers)

October 15, 2023

```
[1]: import pandas as pd
import numpy as np
```

```
[2]: #Reading the CSV file
telco = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")
```

```
[4]: #Displaying column names to identify target variable
columnNames = telco.columns.tolist()
print("Columns: ", columnNames)
```

```
Columns: ['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
'tenure', 'PhoneService', 'MultipleLines', 'InternetService', 'OnlineSecurity',
'OnlineBackup', 'DeviceProtection', 'TechSupport', 'StreamingTV',
'StreamingMovies', 'Contract', 'PaperlessBilling', 'PaymentMethod',
'MonthlyCharges', 'TotalCharges', 'Churn']
```

```
[5]: #1.
#Identifying the target variable
targetVariable = 'Churn'
```

```
[6]: #Counting the number of customers who have churned
telco["Churn"].value_counts()
```

```
[6]: Churn
No      5174
Yes     1869
Name: count, dtype: int64
```

2. The problem we are facing is a binary classification problem in which the datapoint are limited to only two classes(such as yes or no)

```
[8]: #3.
telco.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
```

0	customerID	7043	non-null	object
1	gender	7043	non-null	object
2	SeniorCitizen	7043	non-null	int64
3	Partner	7043	non-null	object
4	Dependents	7043	non-null	object
5	tenure	7043	non-null	int64
6	PhoneService	7043	non-null	object
7	MultipleLines	7043	non-null	object
8	InternetService	7043	non-null	object
9	OnlineSecurity	7043	non-null	object
10	OnlineBackup	7043	non-null	object
11	DeviceProtection	7043	non-null	object
12	TechSupport	7043	non-null	object
13	StreamingTV	7043	non-null	object
14	StreamingMovies	7043	non-null	object
15	Contract	7043	non-null	object
16	PaperlessBilling	7043	non-null	object
17	PaymentMethod	7043	non-null	object
18	MonthlyCharges	7043	non-null	float64
19	TotalCharges	7043	non-null	object
20	Churn	7043	non-null	object

dtypes: float64(1), int64(2), object(18)

memory usage: 1.1+ MB

```
[9]: #4.
      #Checking for missing values
      missingValues = telco.isnull().sum()
      print("Missing values:\n", missingValues)
```

```
Missing values:
  customerID      0
  gender          0
  SeniorCitizen   0
  Partner         0
  Dependents      0
  tenure         0
  PhoneService    0
  MultipleLines   0
  InternetService 0
  OnlineSecurity  0
  OnlineBackup    0
  DeviceProtection 0
  TechSupport     0
  StreamingTV     0
  StreamingMovies 0
  Contract        0
  PaperlessBilling 0
  PaymentMethod   0
```

```
MonthlyCharges    0
TotalCharges       0
Churn              0
dtype: int64
```

4. There are no missing values

```
[10]: #5.
      #Checking for duplicate rows based on all columns
      telcoDuplicate = telco.duplicated(keep=False)

      #Displaying duplicate rows
      print(telco[telcoDuplicate])
```

```
Empty DataFrame
Columns: [customerID, gender, SeniorCitizen, Partner, Dependents, tenure,
PhoneService, MultipleLines, InternetService, OnlineSecurity, OnlineBackup,
DeviceProtection, TechSupport, StreamingTV, StreamingMovies, Contract,
PaperlessBilling, PaymentMethod, MonthlyCharges, TotalCharges, Churn]
Index: []
```

```
[0 rows x 21 columns]
```

5. After executing the above code, it is clear that there are no duplicate rows in the dataset.

```
[11]: #6. Identifying categorical features
      telco.dtypes
```

```
[11]: customerID      object
      gender         object
      SeniorCitizen   int64
      Partner        object
      Dependents      object
      tenure         int64
      PhoneService    object
      MultipleLines    object
      InternetService  object
      OnlineSecurity   object
      OnlineBackup     object
      DeviceProtection object
      TechSupport     object
      StreamingTV     object
      StreamingMovies  object
      Contract        object
      PaperlessBilling object
      PaymentMethod    object
      MonthlyCharges   float64
      TotalCharges     object
      Churn           object
```

dtype: object

```
[12]: #6. I transformed the data using label encoding that assigns a unique numerical
      ↪value to each category. The reason why this
      #code cannot execute is because I have already encoded it previously. The
      ↪verification is among the codes.
      from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['customerID'] = labelEncoder.fit_transform(telco['customerID'])

[13]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['gender'] = labelEncoder.fit_transform(telco['gender'])

[14]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['Partner'] = labelEncoder.fit_transform(telco['Partner'])

[15]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['Dependents'] = labelEncoder.fit_transform(telco['Dependents'])

[16]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['PhoneService'] = labelEncoder.fit_transform(telco['PhoneService'])

[17]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['MultipleLines'] = labelEncoder.fit_transform(telco['MultipleLines'])

[18]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['InternetService'] = labelEncoder.fit_transform(telco['InternetService'])

[19]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['OnlineSecurity'] = labelEncoder.fit_transform(telco['OnlineSecurity'])

[20]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['OnlineBackup'] = labelEncoder.fit_transform(telco['OnlineBackup'])

[21]: from sklearn.preprocessing import LabelEncoder
      labelEncoder = LabelEncoder()
      telco['DeviceProtection'] = labelEncoder.
      ↪fit_transform(telco['DeviceProtection'])
```

```
[22]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['TechSupport'] = labelEncoder.fit_transform(telco['TechSupport'])
```

```
[23]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['StreamingTV'] = labelEncoder.fit_transform(telco['StreamingTV'])
```

```
[24]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['StreamingMovies'] = labelEncoder.fit_transform(telco['StreamingMovies'])
```

```
[25]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['Contract'] = labelEncoder.fit_transform(telco['Contract'])
```

```
[26]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['PaperlessBilling'] = labelEncoder.
    ↪fit_transform(telco['PaperlessBilling'])
```

```
[27]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['PaymentMethod'] = labelEncoder.fit_transform(telco['PaymentMethod'])
```

```
[28]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['TotalCharges'] = labelEncoder.fit_transform(telco['TotalCharges'])
```

```
[29]: from sklearn.preprocessing import LabelEncoder
labelEncoder = LabelEncoder()
telco['Churn'] = labelEncoder.fit_transform(telco['Churn'])
```

```
[30]: telco.dtypes
```

```
[30]: customerID      int32
gender              int32
SeniorCitizen      int64
Partner            int32
Dependents          int32
tenure              int64
PhoneService        int32
MultipleLines       int32
InternetService     int32
OnlineSecurity      int32
OnlineBackup        int32
DeviceProtection    int32
```

```
TechSupport          int32
StreamingTV          int32
StreamingMovies      int32
Contract             int32
PaperlessBilling     int32
PaymentMethod        int32
MonthlyCharges       float64
TotalCharges         int32
Churn                int32
dtype: object
```

```
[31]: telco.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 7043 entries, 0 to 7042
Data columns (total 21 columns):
#   Column                Non-Null Count  Dtype
---  -
0   customerID            7043 non-null   int32
1   gender                7043 non-null   int32
2   SeniorCitizen         7043 non-null   int64
3   Partner               7043 non-null   int32
4   Dependents            7043 non-null   int32
5   tenure                7043 non-null   int64
6   PhoneService          7043 non-null   int32
7   MultipleLines         7043 non-null   int32
8   InternetService       7043 non-null   int32
9   OnlineSecurity        7043 non-null   int32
10  OnlineBackup          7043 non-null   int32
11  DeviceProtection      7043 non-null   int32
12  TechSupport           7043 non-null   int32
13  StreamingTV           7043 non-null   int32
14  StreamingMovies       7043 non-null   int32
15  Contract              7043 non-null   int32
16  PaperlessBilling      7043 non-null   int32
17  PaymentMethod         7043 non-null   int32
18  MonthlyCharges        7043 non-null   float64
19  TotalCharges          7043 non-null   int32
20  Churn                 7043 non-null   int32
dtypes: float64(1), int32(18), int64(2)
memory usage: 660.4 KB
```

```
[33]: #7.
churnDistributions = telco['Churn'].value_counts()
print(churnDistributions)
```

```
Churn
0      5174
```

1 1869

Name: count, dtype: int64

8. An outlier is a datapoint that is separated from other datapoints in a dataset. For instance, if a test was conducted, and majority of the students got good grades like 10,10,10,9,9,10,8,... and so forth but someone got 1, it means that 1 is an outlier because it deviates from the other scores.

```
[39]: #9.
#Checking for outliers
#Q1 = telco['Churn'].quantile(0.25)
#Q3 = telco['Churn'].quantile(0.75)
Q1 = telco.quantile(0.25)
Q3 = telco.quantile(0.75)
IQR = Q3 - Q1

lowerBound = Q1 - 1.5 * IQR
upperBound = Q3 + 1.5 * IQR

#outliers = telco[(telco['Churn'] < lowerBound) | (telco['Churn'] > upperBound)]
#identifying outliers using boolean indexing
outliers = (telco < lowerBound) | (telco > upperBound)
#print(outliers)

#Displaying outliers for each column
outlierRows = telco[outliers.any(axis=1)]
print(outlierRows)
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	\
0	5375	0	0	1	0	1	
3	5535	1	0	0	0	45	
7	4770	0	0	0	0	10	
20	6207	1	1	0	0	1	
27	6119	1	0	1	1	1	
...	...	...	...	...	...	...	
7031	2521	1	1	1	0	55	
7032	4893	1	1	0	0	1	
7036	5504	0	0	0	0	12	
7040	3367	0	0	1	1	11	
7041	5934	1	1	1	0	4	

	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	\
0	0	1	0	0	...	
3	0	1	0	2	...	
7	0	1	0	2	...	
20	0	1	0	0	...	
27	0	1	0	0	...	
...	...	...	...	...	...	
7031	1	2	0	2	...	

7032	1	2	1	0 ...
7036	0	1	0	0 ...
7040	0	1	0	2 ...
7041	1	2	1	0 ...

	DeviceProtection	TechSupport	StreamingTV	StreamingMovies	Contract	\
0	0	0	0	0	0	0
3	2	2	0	0	1	1
7	0	0	0	0	0	0
20	2	0	0	2	0	0
27	0	0	0	0	0	0
...	...	...	...	...	...	...
7031	0	0	0	0	1	1
7032	0	0	0	0	0	0
7036	2	2	2	2	1	1
7040	0	0	0	0	0	0
7041	0	0	0	0	0	0

	PaperlessBilling	PaymentMethod	MonthlyCharges	TotalCharges	Churn
0	1	2	29.85	2505	0
3	0	0	42.30	1400	0
7	0	3	29.75	2609	0
20	1	2	39.65	3340	1
27	0	2	30.20	2592	1
...	...	...	...	...	...
7031	0	1	60.00	2880	0
7032	1	2	75.75	5776	1
7036	0	2	60.65	5741	0
7040	1	2	29.60	2994	0
7041	1	3	74.40	2660	1

[1720 rows x 21 columns]

```
[57]: from sklearn.preprocessing import QuantileTransformer
#9
#Transforming outliers using quantile transformation
quantileTransformer = QuantileTransformer(output_distribution='normal')
telco['ChurnTransformed'] = quantileTransformer.fit_transform(telco[['Churn']])
```

```
-----
ValueError                                Traceback (most recent call last)
Cell In[57], line 5
      2 #9
      3 #Transforming outliers using quantile transformation
      4 quantileTransformer = QuantileTransformer(output_distribution='normal')
----> 5 telco['ChurnTransformed'] = quantileTransformer.
      ↪fit_transform(telco[['Churn']])
```



```

File ~\anaconda3\Lib\site-packages\sklearn\utils\_set_output.py:140, in
↳ _wrap_method_output.<locals>.wrapped(self, X, *args, **kwargs)
    138 @wraps(f)
    139 def wrapped(self, X, *args, **kwargs):
--> 140     data_to_wrap = f(self, X, *args, **kwargs)
    141     if isinstance(data_to_wrap, tuple):
    142         # only wrap the first output for cross decomposition
    143         return_tuple = (
    144             _wrap_data_with_container(method, data_to_wrap[0], X, self)
    145             *data_to_wrap[1:],
    146         )

```

```

File ~\anaconda3\Lib\site-packages\sklearn\base.py:915, in TransformerMixin.
↳ fit_transform(self, X, y, **fit_params)
    911 # non-optimized default implementation; override when a better
    912 # method is possible for a given clustering algorithm
    913 if y is None:
    914     # fit method of arity 1 (unsupervised transformation)
--> 915     return self.fit(X, **fit_params).transform(X)
    916 else:
    917     # fit method of arity 2 (supervised transformation)
    918     return self.fit(X, y, **fit_params).transform(X)

```

```

File ~\anaconda3\Lib\site-packages\sklearn\base.py:1151, in _fit_context.
↳ <locals>.decorator.<locals>.wrapper(estimator, *args, **kwargs)
    1144 estimator._validate_params()
    1146 with config_context(
    1147     skip_parameter_validation=(
    1148         prefer_skip_nested_validation or global_skip_validation
    1149     )
    1150 ):
-> 1151     return fit_method(estimator, *args, **kwargs)

```

```

File ~\anaconda3\Lib\site-packages\sklearn\preprocessing\_data.py:2663, in
↳ QuantileTransformer.fit(self, X, y)
    2656 if self.n_quantiles > self.subsample:
    2657     raise ValueError(
    2658         "The number of quantiles cannot be greater than"
    2659         " the number of samples used. Got {} quantiles"
    2660         " and {} samples.".format(self.n_quantiles, self.subsample)
    2661     )
-> 2663 X = self._check_inputs(X, in_fit=True, copy=False)
    2664 n_samples = X.shape[0]
    2666 if self.n_quantiles > n_samples:

```

```

File ~\anaconda3\Lib\site-packages\sklearn\preprocessing\_data.py:2752, in
↳ QuantileTransformer._check_inputs(self, X, in_fit, accept_sparse_negative,
↳ copy)
    2750 def _check_inputs(self, X, in_fit, accept_sparse_negative=False,
↳ copy=False):
    2751     """Check inputs before fit and transform."""
-> 2752     X = self._validate_data(
    2753         X,
    2754         reset=in_fit,
    2755         accept_sparse="csc",
    2756         copy=copy,
    2757         dtype=FLOAT_DTYPES,
    2758         force_all_finite="allow-nan",
    2759     )
    2760     # we only accept positive sparse matrix when ignore_implicit_zeros
    2761     # false and that we call fit or transform.
    2762     with np.errstate(invalid="ignore"): # hide NaN comparison warnings

```

```

File ~\anaconda3\Lib\site-packages\sklearn\base.py:604, in BaseEstimator.
↳ _validate_data(self, X, y, reset, validate_separately, cast_to_ndarray,
↳ **check_params)
    602         out = X, y
    603     elif not no_val_X and no_val_y:
-> 604         out = check_array(X, input_name="X", **check_params)
    605     elif no_val_X and not no_val_y:
    606         out = _check_y(y, **check_params)

```

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:959, in
↳ check_array(array, accept_sparse, accept_large_sparse, dtype, order, copy,
↳ force_all_finite, ensure_2d, allow_nd, ensure_min_samples,
↳ ensure_min_features, estimator, input_name)
    953         raise ValueError(
    954             "Found array with dim %d. %s expected <= 2."
    955             % (array.ndim, estimator_name)
    956         )
    957     if force_all_finite:
-> 959         _assert_all_finite(
    960             array,
    961             input_name=input_name,
    962             estimator_name=estimator_name,
    963             allow_nan=force_all_finite == "allow-nan",
    964         )
    965     if ensure_min_samples > 0:
    966         n_samples = _num_samples(array)

```

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:124, in
↳ _assert_all_finite(X, allow_nan, msg_dtype, estimator_name, input_name)
    121     if first_pass_isfinite:
    122         return

```

```

--> 124 _assert_all_finite_element_wise(
    125     X,
    126     xp=xp,
    127     allow_nan=allow_nan,
    128     msg_dtype=msg_dtype,
    129     estimator_name=estimator_name,
    130     input_name=input_name,
    131 )

```

File ~\anaconda3\Lib\site-packages\sklearn\utils\validation.py:173, in  
 ↪ \_assert\_all\_finite\_element\_wise(X, xp, allow\_nan, msg\_dtype, estimator\_name,   
 ↪ input\_name)

```

    156 if estimator_name and input_name == "X" and has_nan_error:
    157     # Improve the error message on how to handle missing values in
    158     # scikit-learn.
    159     msg_err += (
    160         f"\n{estimator_name} does not accept missing values"
    161         " encoded as NaN natively. For supervised learning, you might
    ↪ want"
    (... )
    171         "#estimators-that-handle-nan-values"
    172     )
--> 173 raise ValueError(msg_err)

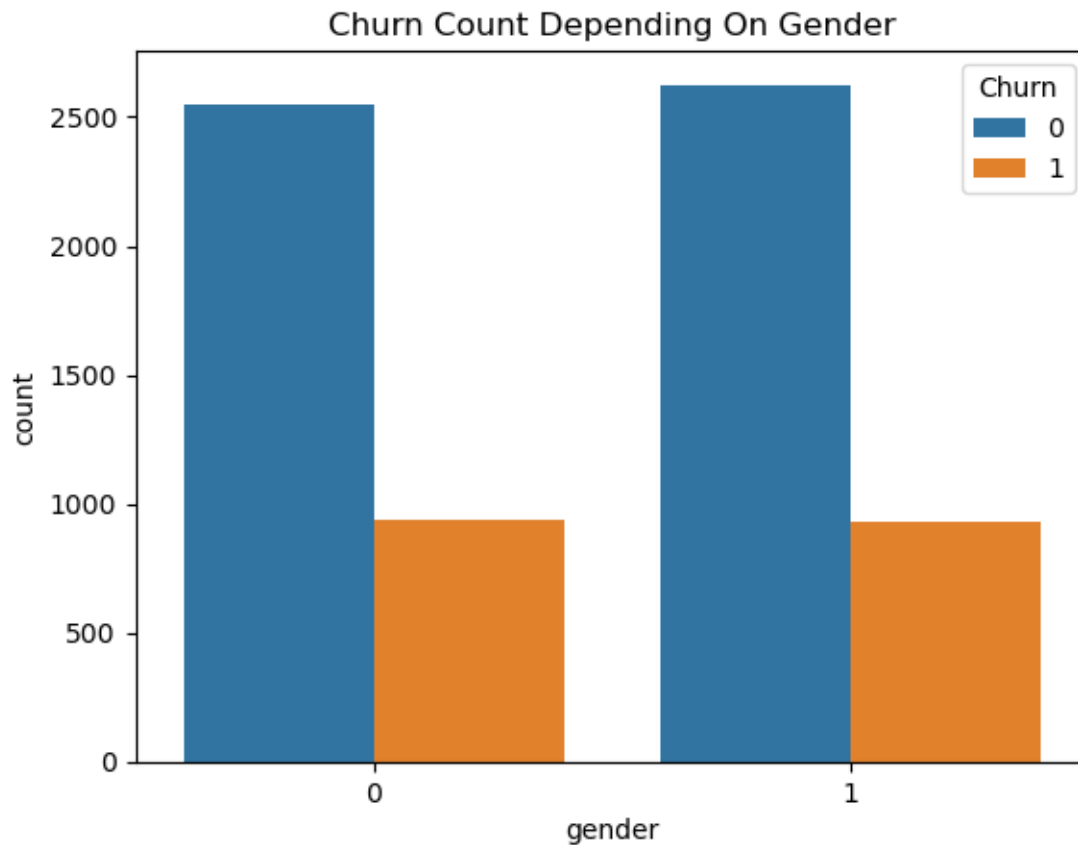
```

**ValueError:** Input X contains infinity or a value too large for dtype('float64')

```


[40]: #10.
import seaborn as sns
import matplotlib.pyplot as plt
sns.countplot(x='gender', hue='Churn', data=telco)
plt.title('Churn Count Depending On Gender')
plt.show()

```

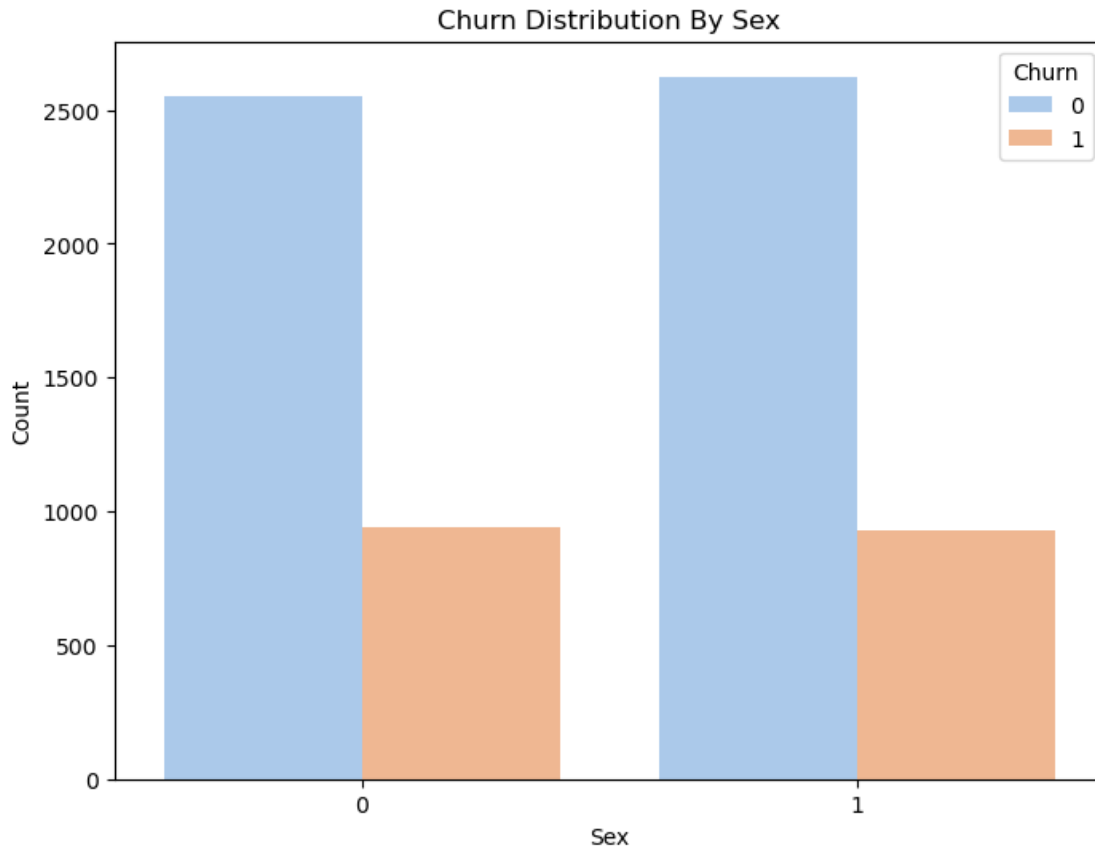


```
[48]: #11.
#Computing percentage of churn customers and active customers
percentageChurn = telco['Churn'].value_counts(normalize=True) * 100
#Using pie chart to visualize
plt.figure(figsize=(6, 6))
plt.pie(percentageChurn, labels=percentageChurn.index, autopct='%1.1f%%',
        ↪startangle=90, colors=['#800080', '#508000'])
plt.title('Churn Customers vs Active Customers')
plt.show()
```

## Churn Customers vs Active Customers



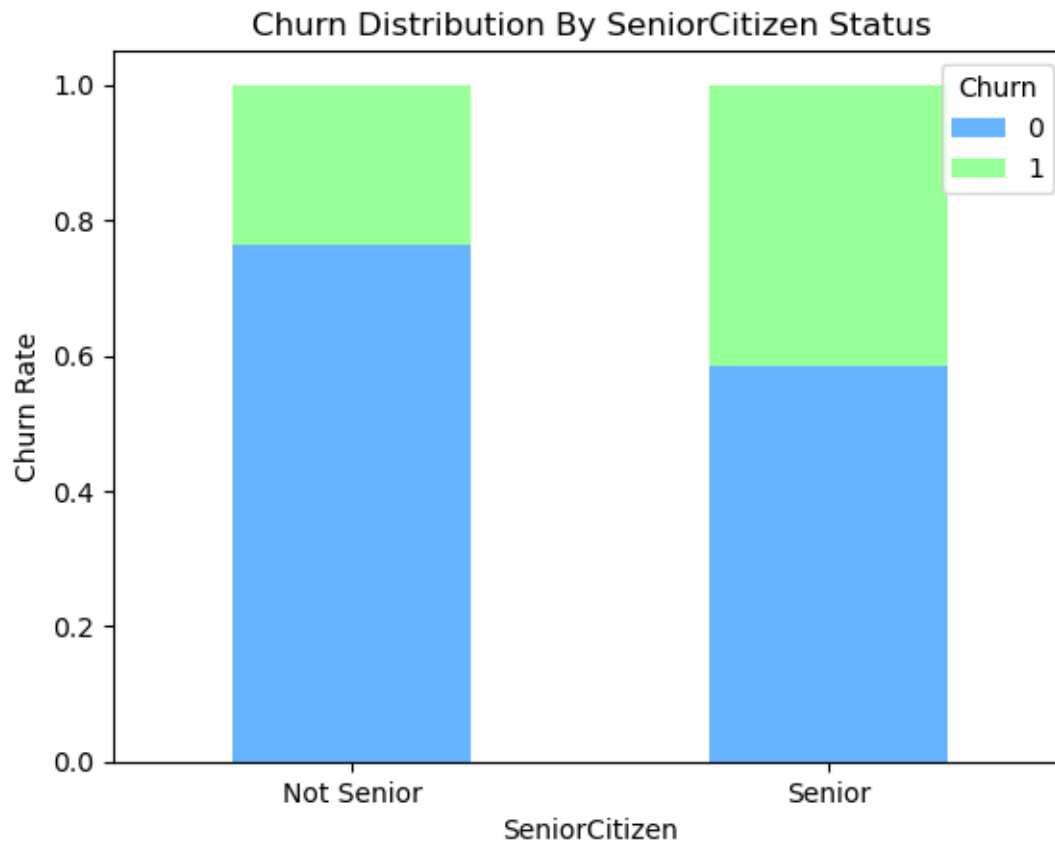
```
[49]: #12. 1st visualization
#Plotting churn distribution based on sex
plt.figure(figsize=(8, 6))
sns.countplot(x='gender', hue='Churn', data=telco, palette='pastel')
plt.title('Churn Distribution By Sex')
plt.xlabel('Sex')
plt.ylabel('Count')
plt.show()
```



```
[51]: #12. 2nd visualization
      #Plotting churn rate based on senior citizen status
      SeniorChurn = telco.groupby('SeniorCitizen')['Churn'].
      ↪value_counts(normalize=True).unstack()

      plt.figure(figsize=(8, 6))
      SeniorChurn.plot(kind='bar', stacked=True, color=['#66b3ff', '#99ff99'])
      plt.title('Churn Distribution By SeniorCitizen Status')
      plt.xlabel('SeniorCitizen')
      plt.ylabel('Churn Rate')
      plt.xticks(ticks=[0, 1], labels=['Not Senior', 'Senior'], rotation=0)
      plt.legend(title='Churn', loc='upper right')
      plt.show()
```

<Figure size 800x600 with 0 Axes>



```
[52]: #12. 3rd visualization
#Visualizing churn rate based on tenure
plt.figure(figsize=(10, 6))
sns.boxplot(x='Churn', hue='tenure', data=telco, palette='pastel')
plt.title('Churn Distribution By Tenure')
plt.xlabel('Churn')
plt.ylabel('Tenure')
plt.show()
```

